

Importing Necessary Libraries

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

import warnings
warnings.filterwarnings(action='ignore')
```

Data Ingestion

```
In [2]: train_df = pd.read_csv('dataset/aps_failure_training_set_processed_8bit.csv')
train_labels = pd.read_csv('dataset/aps_failure_training_set.csv')['class']

test_df = pd.read_csv('dataset/aps_failure_test_set_processed_8bit.csv')
test_labels = pd.read_csv('dataset/aps_failure_test_set.csv')['class']
```

Exploring the Data

```
In [3]: train_df['class'] = train_labels
test_df['class'] = test_labels
```

```
In [4]: train_df.head()
```

```
Out[4]:
```

	class	aa_000	ab_000	ac_000	ad_000	ae_000	af_000	ag_000	ag_001	ag_002	...	ee_002	ee_003	ee_004
0	neg	0.171188	-0.289062	0.992188	-0.007812	-0.046875	-0.054688	-0.007812	-0.03125	-0.054688	...	0.687500	0.515625	0.234375
1	neg	-0.179688	-0.289062	-0.468750	-0.007812	-0.046875	-0.054688	-0.007812	-0.03125	-0.054688	...	-0.023438	-0.062500	-0.132812
2	neg	-0.125000	-0.289062	-0.468750	-0.007812	-0.046875	-0.054688	-0.007812	-0.03125	-0.054688	...	-0.140625	-0.093750	-0.051625
3	neg	-0.406250	-0.289062	-0.468750	-0.007812	-0.046875	-0.007812	-0.007812	-0.03125	-0.054688	...	-0.382812	-0.382812	-0.375000
4	neg	0.007812	-0.289062	-0.468750	-0.007812	-0.046875	-0.054688	-0.007812	-0.03125	-0.054688	...	0.156250	0.031250	-0.031250

```
In [5]: test_df.head()
```

```
Out[5]:
```

	class	aa_000	ab_000	ac_000	ad_000	ae_000	af_000	ag_000	ag_001	ag_002	...	ee_002	ee_003	ee_004
0	neg	-0.406250	-0.289062	-0.46875	-0.007812	-0.046875	-0.054688	-0.007812	-0.03125	-0.054688	...	-0.382812	-0.382812	-0.375000
1	neg	-0.406250	-0.289062	-0.46875	-0.007812	-0.046875	-0.054688	-0.007812	-0.03125	-0.054688	...	-0.382812	-0.382812	-0.375000
2	neg	0.046875	0.554688	-0.46875	-0.007812	-0.046875	-0.054688	-0.007812	-0.03125	-0.054688	...	0.046875	0.312500	-0.000000
3	neg	0.000000	-0.289062	-0.46875	-0.007812	-0.046875	-0.054688	-0.007812	-0.03125	-0.054688	...	0.085938	0.062500	0.031250
4	neg	-0.390625	-0.289062	-0.46875	-0.007812	-0.046875	-0.054688	-0.007812	-0.03125	-0.054688	...	-0.375000	-0.375000	-0.359375

```
In [6]: train_df.describe()
```

```
Out[6]:
```

	class	aa_000	ab_000	ac_000	ad_000	ae_000	af_000	ag_000	ag_001	ag_002	...	ee_002	ee_003	ee_004
count	60000.000000	60000.000000	60000.000000	60000.000000	60000.000000	60000.000000	60000.000000	60000.000000	60000.000000	60000.000000	60000.000000	60000.000000	60000.000000	60000.000000
mean		-0.124611	-0.071121	-0.198529	-0.007737	-0.033483	-0.040633	-0.006584	-0.026241	-0.040699	-0.071121	-0.040699	-0.071121	-0.040699
std		0.367680	0.356812	0.564872	0.004138	0.107086	0.111752	0.030216	0.065200	0.105864	0.180000	0.105864	0.180000	0.180000
min		-0.406250	-0.289062	-0.468750	-0.007812	-0.046875	-0.054688	-0.007812	-0.031250	-0.054688	-0.110000	-0.054688	-0.110000	-0.110000
25%		-0.398438	-0.289062	-0.468750	-0.007812	-0.046875	-0.054688	-0.007812	-0.031250	-0.054688	-0.110000	-0.054688	-0.110000	-0.110000
50%		-0.195312	-0.289062	-0.468750	-0.007812	-0.046875	-0.054688	-0.007812	-0.031250	-0.054688	-0.110000	-0.054688	-0.110000	-0.110000
75%		-0.070312	0.000000	-0.468750	-0.007812	-0.046875	-0.054688	-0.007812	-0.031250	-0.054688	-0.110000	-0.054688	-0.110000	-0.110000
max		0.992188	0.992188	0.992188	0.992188	0.992188	0.992188	0.992188	0.992188	0.992188	0.992188	0.992188	0.992188	0.992188

```
In [7]: train_df.describe(include='all')
```

```
Out[7]:
```

	class	aa_000	ab_000	ac_000	ad_000	ae_000	af_000	ag_000	ag_001	ag_002	...	ee_002	ee_003	ee_004
count	60000	60000.000000	60000.000000	60000.000000	60000.000000	60000.000000	60000.000000	60000.000000	60000.000000	60000.000000	60000.000000	60000.000000	60000.000000	60000.000000
unique	2	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
top	neg	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
freq	59000	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
mean	NaN	-0.124611	-0.071121	-0.198529	-0.007737	-0.033483	-0.040633	-0.006584	-0.026241	-0.040699	...	-0.040699	-0.071121	-0.040699
...
min	NaN	-0.406250	-0.289062	-0.468750	-0.007812	-0.046875	-0.054688	-0.007812	-0.031250	-0.054688	...	-0.054688	-0.110000	-0.110000
25%	NaN	-0.398438	-0.289062	-0.468750	-0.007812	-0.046875	-0.054688	-0.007812	-0.031250	-0.054688	...	-0.054688	-0.110000	-0.110000
50%	NaN	-0.195312	-0.289062	-0.468750	-0.007812	-0.046875	-0.054688	-0.007812	-0.031250	-0.054688	...	-0.054688	-0.110000	-0.110000
75%	NaN	-0.070312	0.000000	-0.468750	-0.007812	-0.046875	-0.054688	-0.007812	-0.031250	-0.054688	...	-0.054688	-0.110000	-0.110000
max	NaN	0.992188	0.992188	0.992188	0.992188	0.992188	0.992188	0.992188	0.992188	0.992188	0.992188	0.992188	0.992188	0.992188

```
In [8]: train_df.isnull().sum()
```

```
Out[8]:
```

	class	aa_000	ab_000	ac_000	ad_000	ae_000	af_000	ag_000	ag_001	ag_002	...	ee_002	ee_003	ee_004
class	0	0	0	0	0	0	0	0	0	0	...	0	0	0
aa_000	0	0	0	0	0	0	0	0	0	0	...	0	0	0
ab_000	0	0	0	0	0	0	0	0	0	0	...	0	0	0
ac_000	0	0	0	0	0	0	0	0	0	0	...	0	0	0
ad_000	0	0	0	0	0	0	0	0	0	0	...	0	0	0
ae_000	0	0	0	0	0	0	0	0	0	0	...	0	0	0
af_000	0	0	0	0	0	0	0	0	0	0	...	0	0	0
ag_000	0	0	0	0	0	0	0	0	0	0	...	0	0	0
ag_001	0	0	0	0	0	0	0	0	0	0	...	0	0	0
ag_002	0	0	0	0	0	0	0	0	0	0	...	0	0	0
...
ee_002	0	0	0	0	0	0	0	0	0	0	...	0	0	0
ee_003	0	0	0	0	0	0	0	0	0	0	...	0	0	0
ee_004	0	0	0	0	0	0	0	0	0	0	...	0	0	0
Length:	171	dtype: int64												

```
In [63]: train_df.corr()
```

```
Out[63]:
```

	aa_000	ab_000	ac_000	ad_000	ae_000	af_000	ag_000	ag_001	ag_002	ag_003	...	ee_002	ee_003	ee_004
aa_000	1.000000	0.121999	0.091271	0.025188	0.110415	0.115131	0.067289	0.185408	0.327367	0.490721	...	0.867362	0.861476	0.881476
ab_000	0.121999	1.000000	0.015372	0.011020	-0.017268	-0.016700	0.016740	0.042054	0.064040	0.076730	...	0.139755	0.124066	0.110000
ac_000	0.091271	0.015372	1.000000	0.005124	0.007997	0.007327	0.021513	0.022585	0.022782	0.020065	...	0.102385	0.104665	0.050000
ad_000	0.025188	0.011020	0.005124	1.000000	-0.001366	-0.001411	0.006440	0.016721	0.017864	0.014560	...	0.028775	0.023424	0.000000
ae_000	0.110415	-0.017268	0.007997	-0.001366	1.000000	0.966834	-0.004659	0.000273	0.037658	0.138005	...	0.037365	0.046288	0.000000
...
af_000	0.115131	-0.016700	0.007327	-0.001411	0.966834	1.000000	0.043951	0.160518	0.354358	0.601018	...	0.559565	0.580086	0.550000
ag_000	0.490721	0.076730	0.020065	0.014560	0.037658	0.601018	1.000000	0.043951	0.354358	0.601018	...	0.559565	0.580086	0.550000
ag_001	0.490721	0.076730	0.020065	0.014560	0.037658	0.601018	0.043951	1.000000	0.043951	0.354358	...	0.559565	0.580086	0.550000
ag_002	0.490721	0.076730	0.020065	0.014560	0.037658	0.601018	0.043951	0.354358	1.000000	0.043951	...	0.559565	0.580086	0.550000
...
ee_002	0.867362	0.861476	0.881476	0.881476	0.881476	0.881476	0.881476	0.881476	0.881476	0.881476	...	1.000000	0.999999	0.999999
ee_003	0.861476	0.881476	0.999999	0.999999	0.999999	0.999999	0.999999	0.999999	0.999999	0.999999	...	0.999999	1.000000	0.999999
ee_004	0.881476	0.999999	0.999999	0.999999	0.999999	0.999999	0.999999	0.999999	0.999999	0.999999	...	0.999999	0.999999	1.000000

```
In [9]: X = train_df.drop('class', axis=1)
y = train_df['class']
```

Splitting into Train & Test Data

```
In [10]: from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test = train_test_split(X,Y,test_size=0.3,random_state=0)
```

```
In [11]: print("Data Shape: {}".format(train_df.shape))
print("Train Shape: {}".format(X_train.shape))
print("Test Shape: {}".format(X_test.shape))
```

```
Data Shape: (60000, 171)
Train Shape: (42000, 170)
Test Shape: (18000, 170)
```

Logistic Regression

```
In [12]: from sklearn.linear_model import LogisticRegression
log_reg = LogisticRegression()
log_reg.fit(X_train,y_train)
```

```
Out[12]: LogisticRegression
LogisticRegression()
```

```
In [13]: print("Test Accuracy: {:.2f}%".format(log_reg.score(X_test, y_test) * 100))
Test Accuracy: 99.12%
```

```
In [14]: test_score = log_reg.score(X_test, y_test)
test_score_per_truck = test_score[X_test.shape[0]]
print("Best model on test set (Cost = $ %.2f):" % test_score)
print("Best model cost per truck on test set (Cost = $ %.2f)" % test_score_per_truck)
```

```
Best model on test set (Cost = $ 0.99):
Best model cost per truck on test set (Cost = $ 0.00)
```

```
In [15]: y_pred_logreg = log_reg.predict(X_test)
```

```
In [16]: from sklearn.metrics import accuracy_score,classification_report
accuracy_score(y_test, y_pred_logreg)
```

```
Out[16]: 0.9911666666666666
```

```
In [17]: print(classification_report(y_test, y_pred_logreg))
```

	precision	recall	f1-score	support
neg	0.99	1.00	1.00	17702
pos	0.82	0.60	0.69	298
accuracy			0.99	18000
macro avg	0.91	0.80	0.84	18000
weighted avg	0.99	0.99	0.99	18000

Decision Tree Classifier

```
In [18]: from sklearn.tree import DecisionTreeClassifier
dt = DecisionTreeClassifier()
dt.fit(X_train,y_train)
```

```
Out[18]: DecisionTreeClassifier
DecisionTreeClassifier()
```

```
In [19]: y_pred_dt = dt.predict(X_test)
accuracy_score(y_test, y_pred_dt)
```

```
Out[19]: 0.9887777777777778
```

```
In [20]: from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred_dt))
```

	precision	recall	f1-score	support
neg	0.99	0.99	0.99	17702
pos	0.66	0.65	0.66	298
accuracy			0.99	18000
macro avg	0.83	0.82	0.83	18000
weighted avg	0.99	0.99	0.99	18000

Support Vector Classifier

```
In [21]: from sklearn.svm import SVC
from sklearn.svm import LinearSVC,SVC
```