

aps_failure_classification__

November 27, 2022

Importing Necessary Libraries

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

import warnings
warnings.filterwarnings(action='ignore')
```

Data Ingestion

```
[2]: train_df = pd.read_csv('dataset/aps_failure_training_set_processed_8bit.csv')
train_labels = pd.read_csv('dataset/aps_failure_training_set.csv')['class']

test_df = pd.read_csv('dataset/aps_failure_test_set_processed_8bit.csv')
test_labels = pd.read_csv('dataset/aps_failure_test_set.csv')['class']
```

Exploring the Data

```
[3]: train_df['class'] = train_labels
test_df['class'] = test_labels
```

```
[4]: train_df.head()
```

```
[4]:  class    aa_000    ab_000    ac_000    ad_000    ae_000    af_000    ag_000  \
0   neg  0.117188 -0.289062  0.992188 -0.007812 -0.046875 -0.054688 -0.007812
1   neg -0.179688 -0.289062 -0.468750 -0.007812 -0.046875 -0.054688 -0.007812
2   neg -0.125000 -0.289062 -0.468750 -0.007812 -0.046875 -0.054688 -0.007812
3   neg -0.406250 -0.289062 -0.468750 -0.007812 -0.046875 -0.007812 -0.007812
4   neg  0.007812 -0.289062 -0.468750 -0.007812 -0.046875 -0.054688 -0.007812

      ag_001    ag_002  ...    ee_002    ee_003    ee_004    ee_005    ee_006  \
0 -0.03125 -0.054688  ...  0.687500  0.515625  0.234375  0.070312  0.007812
1 -0.03125 -0.054688  ... -0.023438 -0.062500 -0.132812 -0.132812 -0.187500
2 -0.03125 -0.054688  ... -0.140625 -0.093750 -0.015625  0.015625 -0.007812
3 -0.03125 -0.054688  ... -0.382812 -0.382812 -0.375000 -0.351562 -0.312500
4 -0.03125 -0.054688  ...  0.156250  0.031250 -0.031250 -0.039062 -0.046875
```

	ee_007	ee_008	ee_009	ef_000	eg_000
0	-0.109375	-0.140625	-0.171875	-0.023438	-0.023438
1	-0.148438	-0.085938	-0.140625	-0.023438	-0.023438
2	-0.109375	-0.093750	-0.164062	-0.023438	-0.023438
3	-0.195312	-0.304688	-0.171875	0.890625	0.992188
4	-0.015625	0.656250	-0.148438	-0.023438	-0.023438

[5 rows x 171 columns]

```
[5]: test_df.head()
```

```
[5]:   class  aa_000  ab_000  ac_000  ad_000  ae_000  af_000  ag_000 \
0   neg -0.406250 -0.289062 -0.46875 -0.007812 -0.046875 -0.054688 -0.007812
1   neg -0.406250 -0.289062 -0.46875 -0.007812 -0.046875 -0.054688 -0.007812
2   neg  0.046875  0.554688 -0.46875 -0.007812 -0.046875 -0.054688 -0.007812
3   neg  0.000000 -0.289062 -0.46875 -0.007812 -0.046875 -0.054688 -0.007812
4   neg -0.390625 -0.289062 -0.46875 -0.007812 -0.046875 -0.054688 -0.007812
```

	ag_001	ag_002	...	ee_002	ee_003	ee_004	ee_005	ee_006	\
0	-0.03125	-0.054688	...	-0.382812	-0.382812	-0.375000	-0.351562	-0.312500	
1	-0.03125	-0.054688	...	-0.382812	-0.382812	-0.375000	-0.351562	-0.312500	
2	-0.03125	-0.054688	...	0.046875	0.312500	-0.000000	-0.109375	0.914062	
3	-0.03125	-0.054688	...	0.085938	0.062500	0.031250	0.085938	0.093750	
4	-0.03125	-0.054688	...	-0.375000	-0.375000	-0.359375	-0.304688	-0.304688	

	ee_007	ee_008	ee_009	ef_000	eg_000
0	-0.195312	-0.304688	-0.171875	-0.023438	-0.023438
1	-0.195312	-0.304688	-0.171875	-0.023438	-0.023438
2	-0.109375	-0.304688	-0.171875	-0.023438	-0.023438
3	-0.078125	0.320312	-0.109375	-0.023438	-0.023438
4	-0.195312	-0.304688	-0.171875	-0.023438	-0.023438

[5 rows x 171 columns]

```
[6]: train_df.describe()
```

```
[6]:
```

	aa_000	ab_000	ac_000	ad_000	ae_000	\
count	60000.000000	60000.000000	60000.000000	60000.000000	60000.000000	
mean	-0.124611	-0.071121	-0.198529	-0.007737	-0.033483	
std	0.367680	0.356812	0.564872	0.004138	0.107086	
min	-0.406250	-0.289062	-0.468750	-0.007812	-0.046875	
25%	-0.398438	-0.289062	-0.468750	-0.007812	-0.046875	
50%	-0.195312	-0.289062	-0.468750	-0.007812	-0.046875	
75%	-0.070312	0.000000	-0.468750	-0.007812	-0.046875	
max	0.992188	0.992188	0.992188	0.992188	0.992188	

	af_000	ag_000	ag_001	ag_002	ag_003	\
--	--------	--------	--------	--------	--------	---

count	60000.000000	60000.000000	60000.000000	60000.000000	60000.000000
mean	-0.040633	-0.006584	-0.026241	-0.040699	-0.074768
std	0.111752	0.032016	0.065200	0.105864	0.186822
min	-0.054688	-0.007812	-0.031250	-0.054688	-0.117188
25%	-0.054688	-0.007812	-0.031250	-0.054688	-0.117188
50%	-0.054688	-0.007812	-0.031250	-0.054688	-0.117188
75%	-0.054688	-0.007812	-0.031250	-0.054688	-0.117188
max	0.992188	0.992188	0.992188	0.992188	0.992188

	...	ee_002	ee_003	ee_004	ee_005	\
count	...	60000.000000	60000.000000	60000.000000	60000.000000	
mean	...	-0.104808	-0.098734	-0.094976	-0.089227	
std	...	0.356547	0.362066	0.363148	0.336121	
min	...	-0.382812	-0.382812	-0.382812	-0.351562	
25%	...	-0.382812	-0.382812	-0.375000	-0.343750	
50%	...	-0.179688	-0.179688	-0.195312	-0.179688	
75%	...	-0.007812	0.015625	0.015625	0.007812	
max	...	0.992188	0.992188	0.992188	0.992188	

		ee_006	ee_007	ee_008	ee_009	ef_000	\
count	60000.000000	60000.000000	60000.000000	60000.000000	60000.000000	60000.000000	
mean	-0.103374	-0.088961	-0.084540	-0.067471	-0.020035		
std	0.320314	0.237613	0.363893	0.261009	0.051907		
min	-0.312500	-0.195312	-0.304688	-0.171875	-0.023438		
25%	-0.312500	-0.195312	-0.304688	-0.171875	-0.023438		
50%	-0.226562	-0.171875	-0.296875	-0.171875	-0.023438		
75%	-0.054688	-0.101562	0.000000	-0.132812	-0.023438		
max	0.992188	0.992188	0.992188	0.992188	0.992188		

	eg_000
count	60000.000000
mean	-0.018417
std	0.061751
min	-0.023438
25%	-0.023438
50%	-0.023438
75%	-0.023438
max	0.992188

[8 rows x 170 columns]

```
[7]: train_df.describe(include='all')
```

	class	aa_000	ab_000	ac_000	ad_000	\
count	60000	60000.000000	60000.000000	60000.000000	60000.000000	
unique	2	NaN	NaN	NaN	NaN	
top	neg	NaN	NaN	NaN	NaN	

freq	59000	NaN	NaN	NaN	NaN
mean	NaN	-0.124611	-0.071121	-0.198529	-0.007737
std	NaN	0.367680	0.356812	0.564872	0.004138
min	NaN	-0.406250	-0.289062	-0.468750	-0.007812
25%	NaN	-0.398438	-0.289062	-0.468750	-0.007812
50%	NaN	-0.195312	-0.289062	-0.468750	-0.007812
75%	NaN	-0.070312	0.000000	-0.468750	-0.007812
max	NaN	0.992188	0.992188	0.992188	0.992188

	ae_000	af_000	ag_000	ag_001	ag_002 \
count	60000.000000	60000.000000	60000.000000	60000.000000	60000.000000
unique	NaN	NaN	NaN	NaN	NaN
top	NaN	NaN	NaN	NaN	NaN
freq	NaN	NaN	NaN	NaN	NaN
mean	-0.033483	-0.040633	-0.006584	-0.026241	-0.040699
std	0.107086	0.111752	0.032016	0.065200	0.105864
min	-0.046875	-0.054688	-0.007812	-0.031250	-0.054688
25%	-0.046875	-0.054688	-0.007812	-0.031250	-0.054688
50%	-0.046875	-0.054688	-0.007812	-0.031250	-0.054688
75%	-0.046875	-0.054688	-0.007812	-0.031250	-0.054688
max	0.992188	0.992188	0.992188	0.992188	0.992188

	...	ee_002	ee_003	ee_004	ee_005 \
count	...	60000.000000	60000.000000	60000.000000	60000.000000
unique	...	NaN	NaN	NaN	NaN
top	...	NaN	NaN	NaN	NaN
freq	...	NaN	NaN	NaN	NaN
mean	...	-0.104808	-0.098734	-0.094976	-0.089227
std	...	0.356547	0.362066	0.363148	0.336121
min	...	-0.382812	-0.382812	-0.382812	-0.351562
25%	...	-0.382812	-0.382812	-0.375000	-0.343750
50%	...	-0.179688	-0.179688	-0.195312	-0.179688
75%	...	-0.007812	0.015625	0.015625	0.007812
max	...	0.992188	0.992188	0.992188	0.992188

	ee_006	ee_007	ee_008	ee_009	ef_000 \
count	60000.000000	60000.000000	60000.000000	60000.000000	60000.000000
unique	NaN	NaN	NaN	NaN	NaN
top	NaN	NaN	NaN	NaN	NaN
freq	NaN	NaN	NaN	NaN	NaN
mean	-0.103374	-0.088961	-0.084540	-0.067471	-0.020035
std	0.320314	0.237613	0.363893	0.261009	0.051907
min	-0.312500	-0.195312	-0.304688	-0.171875	-0.023438
25%	-0.312500	-0.195312	-0.304688	-0.171875	-0.023438
50%	-0.226562	-0.171875	-0.296875	-0.171875	-0.023438
75%	-0.054688	-0.101562	0.000000	-0.132812	-0.023438
max	0.992188	0.992188	0.992188	0.992188	0.992188

```

          eg_000
count    60000.000000
unique           NaN
top           NaN
freq           NaN
mean      -0.018417
std        0.061751
min       -0.023438
25%       -0.023438
50%       -0.023438
75%       -0.023438
max        0.992188

```

[11 rows x 171 columns]

Checking the Null Values

```
[8]: train_df.isnull().sum()
```

```

[8]: class      0
     aa_000     0
     ab_000     0
     ac_000     0
     ad_000     0
     ..
     ee_007     0
     ee_008     0
     ee_009     0
     ef_000     0
     eg_000     0
     Length: 171, dtype: int64

```

```
[63]: train_df.corr()
```

```

[63]:
          aa_000  ab_000  ac_000  ad_000  ae_000  af_000  ag_000  \
aa_000  1.000000  0.121999  0.091271  0.025188  0.110415  0.115131  0.067289
ab_000  0.121999  1.000000  0.015372  0.011020 -0.017268 -0.016700  0.016740
ac_000  0.091271  0.015372  1.000000  0.005124  0.007997  0.007327  0.021513
ad_000  0.025188  0.011020  0.005124  1.000000 -0.001366 -0.001411  0.006440
ae_000  0.110415 -0.017268  0.007997 -0.001366  1.000000  0.966834 -0.004659
...
ee_007  0.714641  0.103732  0.066857  0.015753  0.124427  0.128893  0.043951
ee_008  0.488417  0.034100  0.072585  0.004279 -0.058660 -0.057943  0.017027
ee_009  0.302424  0.021730  0.047044  0.003246 -0.049432 -0.049671 -0.000373
ef_000  0.048662  0.011490  0.011207  0.000651  0.221157  0.218371  0.004524
eg_000  0.025319  0.004143  0.017086  0.000069  0.204498  0.202610  0.002599

```

	ag_001	ag_002	ag_003	...	ee_002	ee_003	ee_004	\
aa_000	0.185408	0.327367	0.490721	...	0.867362	0.861476	0.839020	
ab_000	0.042054	0.064040	0.076730	...	0.139755	0.124066	0.115553	
ac_000	0.022585	0.022782	0.020065	...	0.102385	0.104665	0.097855	
ad_000	0.016721	0.017864	0.014560	...	0.028775	0.023424	0.020111	
ae_000	0.000273	0.037658	0.138005	...	0.037365	0.046288	0.029458	
...	
ee_007	0.160518	0.354358	0.601018	...	0.595695	0.580086	0.556076	
ee_008	0.021267	0.024793	0.018597	...	0.480937	0.482618	0.449890	
ee_009	-0.020576	-0.045709	-0.081193	...	0.295910	0.292490	0.266447	
ef_000	0.019810	0.047570	0.080130	...	0.019257	0.017986	0.010572	
eg_000	0.007825	0.030623	0.055827	...	-0.002936	-0.002632	-0.011163	

	ee_005	ee_006	ee_007	ee_008	ee_009	ef_000	eg_000
aa_000	0.839522	0.823592	0.714641	0.488417	0.302424	0.048662	0.025319
ab_000	0.118397	0.103143	0.103732	0.034100	0.021730	0.011490	0.004143
ac_000	0.097390	0.083394	0.066857	0.072585	0.047044	0.011207	0.017086
ad_000	0.016272	0.013391	0.015753	0.004279	0.003246	0.000651	0.000069
ae_000	0.050052	0.177118	0.124427	-0.058660	-0.049432	0.221157	0.204498
...
ee_007	0.600210	0.766848	1.000000	0.374890	0.187208	0.036857	0.017779
ee_008	0.473259	0.483717	0.374890	1.000000	0.780811	-0.025096	-0.035292
ee_009	0.272294	0.268450	0.187208	0.780811	1.000000	-0.018435	-0.024006
ef_000	0.021801	0.046217	0.036857	-0.025096	-0.018435	1.000000	0.632963
eg_000	0.004524	0.034035	0.017779	-0.035292	-0.024006	0.632963	1.000000

[170 rows x 170 columns]

```
[9]: X = train_df.drop('class', axis=1)
      Y = train_df['class']
```

Splitting into Train & Test Data

```
[10]: from sklearn.model_selection import train_test_split
      X_train,X_test,y_train,y_test = train_test_split(X,Y,test_size=0.
      ↪3,random_state=0)
```

```
[11]: print("Data Shape: {}".format(train_df.shape))
      print("Train Shape: {}".format(X_train.shape))
      print("Test Shape: {}".format(X_test.shape))
```

Data Shape: (60000, 171)

Train Shape: (42000, 170)

Test Shape: (18000, 170)

Logistic Regression

```
[12]: from sklearn.linear_model import LogisticRegression
      log_reg = LogisticRegression()
```

```
log_reg.fit(X_train, y_train)
```

```
[12]: LogisticRegression()
```

```
[13]: print("Test Accuracy: {:.2f}%".format(log_reg.score(X_test, y_test) * 100))
```

Test Accuracy: 99.12%

```
[14]: test_score = log_reg.score(X_test, y_test)
test_score_per_truck = test_score/X_test.shape[0]

print("Best model on test set (Cost = $ %.2f):" % test_score)
print("Best model cost per truck on test set (Cost = $ %.2f)" %_
↪test_score_per_truck)
```

Best model on test set (Cost = \$ 0.99):

Best model cost per truck on test set (Cost = \$ 0.00)

```
[15]: y_pred_logreg = log_reg.predict(X_test)
```

```
[16]: from sklearn.metrics import accuracy_score,classification_report
accuracy_score(y_test, y_pred_logreg)
```

```
[16]: 0.9911666666666666
```

```
[17]: print(classification_report(y_test, y_pred_logreg))
```

	precision	recall	f1-score	support
neg	0.99	1.00	1.00	17702
pos	0.82	0.60	0.69	298
accuracy			0.99	18000
macro avg	0.91	0.80	0.84	18000
weighted avg	0.99	0.99	0.99	18000

Decision Tree Classifier

```
[18]: from sklearn.tree import DecisionTreeClassifier
dt = DecisionTreeClassifier()
dt.fit(X_train,y_train)
```

```
[18]: DecisionTreeClassifier()
```

```
[19]: y_pred_dt = dt.predict(X_test)
accuracy_score(y_test, y_pred_dt)
```

```
[19]: 0.9887777777777778
```

```
[20]: from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred_dt))
```

	precision	recall	f1-score	support
neg	0.99	0.99	0.99	17702
pos	0.66	0.65	0.66	298
accuracy			0.99	18000
macro avg	0.83	0.82	0.83	18000
weighted avg	0.99	0.99	0.99	18000

Support Vector Classifier

```
[21]: from sklearn.svm import SVC
from sklearn.svm import LinearSVC,SVC
from sklearn.pipeline import make_pipeline
clf = SVC(C = 1.2, gamma = 0.9, kernel= 'rbf')
```

```
[22]: clf = clf.fit(X_train,y_train)
y_pred_svc = clf.predict(X_test)
```

```
[23]: from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred_svc))
```

	precision	recall	f1-score	support
neg	0.98	1.00	0.99	17702
pos	1.00	0.04	0.07	298
accuracy			0.98	18000
macro avg	0.99	0.52	0.53	18000
weighted avg	0.98	0.98	0.98	18000

```
[24]: accuracy_score(y_test, y_pred_svc)
```

```
[24]: 0.9840555555555556
```

Ensemble Methods

```
[25]: from sklearn.tree import ExtraTreeClassifier
from sklearn.ensemble import BaggingClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import VotingClassifier
```

```
[26]: extra_model = ExtraTreeClassifier()
extra_model.fit(X,Y)
feature_imp = extra_model.feature_importances_
```



```
[27]: for index, val in enumerate(feature_imp):  
      print(index, round((val * 100), 2))
```

```
0 0.99  
1 0.27  
2 0.7  
3 0.02  
4 0.04  
5 0.03  
6 0.06  
7 5.36  
8 1.15  
9 13.38  
10 0.12  
11 0.12  
12 0.0  
13 0.4  
14 0.17  
15 0.28  
16 5.71  
17 1.01  
18 0.48  
19 0.23  
20 0.55  
21 0.77  
22 0.24  
23 0.04  
24 3.06  
25 0.37  
26 0.19  
27 0.0  
28 0.23  
29 0.09  
30 0.33  
31 0.12  
32 0.01  
33 0.51  
34 2.78  
35 0.43  
36 0.4  
37 3.8  
38 0.8  
39 1.13  
40 2.3  
41 0.28  
42 0.71  
43 0.22  
44 0.59
```

45 0.56
46 0.54
47 0.37
48 0.2
49 0.08
50 0.15
51 0.39
52 0.31
53 0.56
54 0.15
55 0.11
56 0.19
57 0.44
58 0.27
59 0.34
60 0.57
61 0.54
62 0.32
63 0.52
64 0.66
65 0.59
66 0.27
67 0.26
68 0.14
69 0.36
70 1.82
71 0.34
72 0.85
73 0.23
74 0.41
75 0.72
76 0.45
77 0.72
78 0.78
79 0.72
80 0.11
81 0.05
82 0.06
83 0.11
84 0.17
85 0.64
86 0.45
87 0.27
88 0.15
89 0.17
90 0.3
91 0.17
92 0.09

93 0.09
94 0.06
95 0.29
96 0.09
97 5.29
98 0.71
99 1.72
100 0.4
101 1.16
102 0.73
103 0.26
104 0.07
105 0.52
106 0.38
107 0.52
108 0.31
109 0.07
110 0.36
111 0.29
112 0.23
113 0.14
114 0.74
115 1.02
116 0.24
117 0.33
118 0.17
119 0.87
120 0.3
121 0.24
122 0.0
123 0.87
124 0.29
125 0.24
126 0.27
127 0.27
128 0.03
129 0.23
130 0.3
131 0.41
132 0.32
133 0.32
134 0.18
135 0.14
136 0.12
137 0.21
138 0.0
139 0.08
140 0.0

```

141 0.03
142 0.22
143 0.18
144 0.33
145 0.07
146 0.13
147 0.2
148 0.29
149 0.19
150 0.61
151 0.33
152 1.05
153 0.23
154 0.0
155 0.51
156 0.32
157 0.37
158 0.48
159 0.23
160 0.03
161 0.15
162 0.21
163 0.21
164 0.25
165 1.77
166 0.18
167 0.45
168 0.0
169 0.05

```

Random Forest Classifier

```
[28]: rf = RandomForestClassifier()
      rf.fit(X_train,y_train)
```

```
[28]: RandomForestClassifier()
```

```
[29]: y_pred_rf = rf.predict(X_test)
      accuracy_score(y_test,y_pred_rf)
```

```
[29]: 0.9923333333333333
```

```
[30]: from sklearn.metrics import classification_report
      print(classification_report(y_test, y_pred_rf))
```

	precision	recall	f1-score	support
neg	0.99	1.00	1.00	17702
pos	0.93	0.58	0.71	298

accuracy			0.99	18000
macro avg	0.96	0.79	0.85	18000
weighted avg	0.99	0.99	0.99	18000

Bagging Classifier

```
[31]: cls = BaggingClassifier(rf, random_state=0).fit(X_train, y_train)
      cls.score(X_test, y_test)
```

```
[31]: 0.99144444444444445
```

```
[32]: y_pred_clf = cls.predict(X_test)
```

```
[33]: accuracy_score(y_test, y_pred_clf)
```

```
[33]: 0.99144444444444445
```

```
[34]: from sklearn.metrics import classification_report
      print(classification_report(y_test, y_pred_clf))
```

	precision	recall	f1-score	support
neg	0.99	1.00	1.00	17702
pos	0.93	0.52	0.67	298
accuracy			0.99	18000
macro avg	0.96	0.76	0.83	18000
weighted avg	0.99	0.99	0.99	18000

Extra Tree Classifier

```
[35]: extra_tree = ExtraTreeClassifier(random_state=0)
      cls_extra = BaggingClassifier(extra_tree, random_state=0).fit(X_train, y_train)
      cls_extra.score(X_test, y_test)
```

```
[35]: 0.9908888888888889
```

Voting Classifier

```
[36]: clf1 = ExtraTreeClassifier(random_state=0)
      clf2 = RandomForestClassifier()
      clf3 = DecisionTreeClassifier()
      clf4 = LogisticRegression()
      clf5 = SVC(C = 1.2, gamma = 0.9, kernel= 'rbf')
```

```
[37]: eclf1 = VotingClassifier(estimators=[
      ('ETC', clf1), ('RF', clf2), ('DT', clf3), ('Log_Reg', clf4), ('SVC',
      ↪clf5), ('Bagging', cls)], voting='hard')
```

```
[38]: eclf1 = eclf1.fit(X, Y)
      print(eclf1.predict(X))
```

```
['neg' 'neg' 'neg' ... 'neg' 'neg' 'neg']
```

Boosting Algorithms

Adaboost Classifier

```
[39]: from sklearn.ensemble import AdaBoostClassifier
      ada_model = AdaBoostClassifier()
```

```
[41]: ada_model.fit(X_train,y_train)
```

```
[41]: AdaBoostClassifier()
```

```
[42]: y_pred_ada = ada_model.predict(X_test)
```

```
[43]: accuracy_score(y_test,y_pred_ada)
```

```
[43]: 0.9904444444444445
```

```
[44]: from sklearn.metrics import classification_report
      print(classification_report(y_test, y_pred_ada))
```

	precision	recall	f1-score	support
neg	0.99	1.00	1.00	17702
pos	0.76	0.61	0.68	298
accuracy			0.99	18000
macro avg	0.88	0.81	0.84	18000
weighted avg	0.99	0.99	0.99	18000

Gradient Boosting Algorithm

```
[45]: from sklearn.ensemble import GradientBoostingClassifier
      GB_model = GradientBoostingClassifier(n_estimators=100, learning_rate=1.
      ↪0,max_depth=1, random_state=0)
```

```
[46]: GB_model.fit(X_train,y_train)
```

```
[46]: GradientBoostingClassifier(learning_rate=1.0, max_depth=1, random_state=0)
```

```
[47]: y_pred_gb = ada_model.predict(X_test)
      accuracy_score(y_test,y_pred_gb)
```

```
[47]: 0.9904444444444445
```

```
[51]: from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred_gb))
```

	precision	recall	f1-score	support
neg	0.99	1.00	1.00	17702
pos	0.76	0.61	0.68	298
accuracy			0.99	18000
macro avg	0.88	0.81	0.84	18000
weighted avg	0.99	0.99	0.99	18000

```
[ ]:
```